

Weighted Model Counting Without Parameter Variables

No Author Given

No Institute Given

Abstract. The abstract should briefly summarize the contents of the paper in 150–250 words.

Keywords: First keyword · Second keyword · Another keyword.

1 Introduction

Related Work. Many WMC inference algorithms such as *Ace*¹, *c2d* [10], and *miniC2D* [14] work by compilation to tractable representations such as arithmetic circuits, deterministic, decomposable negation normal form [8], and sentential decision diagrams (SDDs) [11]. Some attempts have previously been made to skip this intermediate stage of representing a Bayesian network as a WMC instance and compile directly to a more convenient representation. Specifically, direct compilation from Bayesian networks to SDDs [7] and from structured Bayesian networks (i.e., a generalisation of Bayesian networks) to probabilistic SDDs [17] have been considered. To the best of the authors’ knowledge, neither compilation approach has a publicly available implementation.

Our work is similar in that it introduces a new representation for computational problems that is based on pseudo-Boolean functions—pseudo-Boolean projection (PBP). We formally show that every WMC problem instance has a corresponding PBP instance. We argue that this new format is more appropriate for WMC solvers whose execution is build on pseudo-Boolean function manipulation, i.e., *ADDMC* [12] and *DPMC* [13]. Finally, we show how most methods of encoding probabilistic inference on Bayesian networks to WMC can be adapted to the PBP format. This transformation significantly improves inference speed, outperforming all other considered algorithm-encoding combinations.

2 Pseudo-Boolean Functions

Notation. For any propositional formula ϕ over a set of variables X and $p, q \in \mathbb{R}$, let $[\phi]_q^p: 2^X \rightarrow \mathbb{R}$ be a pseudo-Boolean function defined as

$$[\phi]_q^p(Y) := \begin{cases} p & \text{if } Y \models \phi \\ q & \text{otherwise} \end{cases}$$

for any $Y \subseteq X$.

¹ <http://reasoning.cs.ucla.edu/ace/>

Definition 1 (Operations). Let $f, g: 2^X \rightarrow \mathbb{R}$ be pseudo-Boolean functions, $x, y \in X$, $Y = \{y_i\}_{i=1}^n \subseteq X$, and $r \in \mathbb{R}$. Operations such as addition and multiplication are defined pointwise as

$$(f + g)(Y) := f(Y) + g(Y), \quad \text{and} \quad (f \cdot g)(Y) := f(Y) \cdot g(Y).$$

Note that this means that binary operations on pseudo-Boolean functions inherit properties such as associativity and commutativity. By not distinguishing between a real number and a pseudo-Boolean function that always returns that number, we can use the same definitions to define scalar operations as

$$(r + f)(Y) = r + f(Y), \quad \text{and} \quad (r \cdot f)(Y) = r \cdot f(Y).$$

Restrictions $f|_{x=0}, f|_{x=1}: 2^X \rightarrow \mathbb{R}$ of f are defined as

$$f|_{x=0}(Y) := f(Y \setminus \{x\}), \quad \text{and} \quad f|_{x=1}(Y) := f(Y \cup \{x\})$$

for all $Y \subseteq X$.

Projection \exists_x is an endomorphism $\exists_x: \mathbb{R}^{2^X} \rightarrow \mathbb{R}^{2^X}$ defined as

$$\exists_x f := f|_{x=1} + f|_{x=0}.$$

Since projection is commutative (i.e., $\exists_x \exists_y f = \exists_y \exists_x f$) [12, 13], we can define $\exists_Y: \mathbb{R}^{2^X} \rightarrow \mathbb{R}^{2^X}$ as $\exists_Y := \exists_{y_1} \exists_{y_2} \dots \exists_{y_n}$. Throughout the paper, projection is assumed to have the lowest precedence (e.g., $\exists_x fg = \exists_x (fg)$).

NOTE: We list some properties of the operations on pseudo-Boolean functions discussed in this section that can be conveniently represented using our syntax. The proofs of all these properties follow directly from the definitions.

Proposition 1 (Basic Properties). For any propositional formulas ϕ and ψ , and $a, b, c, d \in \mathbb{R}$,

- $[\phi]_b^a = [\neg\phi]_a^b$;
- $c + [\phi]_b^a = [\phi]_{b+c}^{a+c}$;
- $c \cdot [\phi]_b^a = [\phi]_{bc}^{ac}$;
- $[\phi]_b^a \cdot [\phi]_d^c = [\phi]_{bd}^{ac}$;
- $[\phi]_0^1 \cdot [\psi]_0^1 = [\phi \wedge \psi]_0^1$.

And for any pair of pseudo-Boolean functions $f, g: 2^X \rightarrow \mathbb{R}$ and $x \in X$, $(fg)|_{x=i} = f|_{x=i} \cdot g|_{x=i}$ for $i = 0, 1$.

Notes.

- Weights on literals other than the positive literals that correspond to variables in X_P are redundant as they either are equal to one or duplicate an already-defined weight.
- Mention that formulas, clauses, and models are all treated as sets.

3 Weighted Model Counting

NOTE: Our definition of WMC is largely based on the standard definition [5], but explicitly partitions variables into indicator and parameter variables.

Definition 2 (WMC Instance). A WMC instance is a tuple (ϕ, X_I, X_P, w) , where X_I is the set of indicator variables, X_P is the set of parameter variables (with $X_I \cap X_P = \emptyset$), ϕ is a propositional formula in CNF over $X_I \cup X_P$, and $w: X_I \cup X_P \cup \{\neg x \mid x \in X_I \cup X_P\} \rightarrow \mathbb{R}$ is the weight function. The answer of the instance is $\sum_{Y \models \phi} \prod_{Y \models l} w(l)$.

Definition 3 (PBP Instance). A PBP instance is a tuple (F, X, ω) , where X is the set of variables, F is a set of pseudo-Boolean functions $2^X \rightarrow \mathbb{R}$, and $\omega \in \mathbb{R}$ is the scaling factor.² The answer of the instance is $\omega \cdot \left(\exists_X \prod_{f \in F} f \right) (\emptyset)$.

NOTE: Adding scaling factor ω to the definition allows us to remove clauses that consist entirely of a single parameter variable. The idea of extracting some of the structure of the WMC instance into an external multiplicative factor originates in one of the Bayesian network encodings [2].

3.1 Bayesian Network Encodings

TODO

- Define the terminology surrounding BNs: variables, values, directed probabilistic graphical model, DAG, binary = two values per variable.
- Reuse this term where I'm referring to two values.
- Give a brief description of each encoding.
- Talk about the parameter/indicator variable distinction.

WMC is a well-established technique for Bayesian network inference, particularly effective on networks where most variables have only a few possible values [9]. Many ways of encoding a Bayesian network into a WMC instance have been proposed. We will refer to them based on initials of the authors and the year of publication: **d02** [9], **sbk05** [16], **cd05** [3], **cd06** [4], **bk1m16** [2].

Remark 1. Strictly speaking, encodings such as **cd05** and **cd06** are *not* WMC encodings. Instead, they encode Bayesian network inference into instances of the *minimum-cardinality* WMC problem, where the answer is defined to be $\sum_{Y \models \phi, |Y|=k} \prod_{Y \models l} w(l)$, where $k = \min_{Y \models \phi, Y \neq \emptyset} |Y|$, if k exists, otherwise the answer is zero. This additional condition on model cardinality becomes necessary because these encodings eliminate clauses of the form $p \Rightarrow i$, where $p \in X_P$ is a

² We will further restrict F to only contain functions that have exactly two values. While this restriction is not necessary from a computational point of view, such functions can be conveniently represented both in text files and using the syntax proposed in this paper.

parameter variable, and $i \in X_I$ is an indicator variable. Nonetheless, our transformation algorithm still works on such encodings, although the experimental results are discouraging because they use approximately twice as many indicator variables. For instance, each binary variable of a Bayesian network is encoded using two indicator variables while one would suffice.

4 Parameter Variable Elimination

Notes.

- Let X_P be the set of parameter variable and X_I be the set of indicator variables.
- Parameter variables are either taken from the LMAP file (for encodings produced by Ace) or assumed to be the variables that have both weights equal to 1.
- If a parameter variable in a clause is ‘negated’, we can ignore the clause. We assume that there are no clauses with more than one instance of parameter variables.
- The second **foreach** loop can be performed in constant time by representing ϕ' as a list and assuming that the two ‘clauses’ are adjacent in that list (and incorporating it into the first loop).
- The d map is constructed in $\mathcal{O}(|X_P| \log |X_P|)$ time (we want to use a data structure based on binary search trees rather than hashing).
- **rename** can be implemented in $\mathcal{O}(\log |X_P|)$ time.
- This may look like preprocessing, but all the transformations are local and thus can be incorporated into an encoding algorithm with no slowdown. In fact, if anything, the resulting algorithm would be slightly faster, as it would have less data to output.
- We do this as a follow-up to Ace compilation because the implementations of the encodings are all closed-source, and we don’t want to re-invent (a subpar version of) the wheel.
- No obvious way to do this for sbk05 because the roles of indicator and parameter (i.e., ‘chance’) variables overlap.
- This just says that p is equivalent to a conjunction.
- The first group of conditions applies to d02, while the second group applies to bklm16.
- For cd05 and cd06, condition 2*bi* should be replaced with $w(\neg p) = 1$.
- Benefits of having this proof in the paper:
 - It puts all encodings on a common ground.
 - It illustrates the convenience of our notation for reasoning about (certain types of) pseudo-Boolean functions.
 - It’s too big and too important to be left for the appendix.

Algorithm 1: WMC instance transformation

Data: WMC instance (ϕ, X_I, X_P, w)
Result: PBP instance (F, X, ω)

```

1   $F \leftarrow \emptyset;$ 
2   $\omega \leftarrow 1;$ 
3  let  $d: X_P \rightarrow \mathbb{N}$  be defined as  $p \mapsto |\{o \in X_P \mid o \leq p\}|;$ 
4  foreach clause  $c \in \phi$  do
5      if  $c \cap X_P = \{p\}$  for some  $p$  and  $w(p) \neq 1$  then
6          if  $|c| = 1$  then
7               $\omega \leftarrow \omega \times w(p);$ 
8          else
9               $F \leftarrow F \cup \left\{ \left[ \bigwedge_{l \in c \setminus \{p\}} \neg l \right]_1^{w(p)} \right\};$ 
10         else if  $\{p \mid \neg p \in c\} \cap X_P = \emptyset$  then
11              $F \leftarrow F \cup \{[c]_0^1\};$ 
12 foreach indicator variable  $v \in X_I$  do
13     if  $\{[v]_1^p, [\neg v]_1^q\} \subseteq F$  for some  $p$  and  $q$  then
14          $F \leftarrow F \setminus \{[v]_1^p, [\neg v]_1^q\} \cup \{[v]_q^p\};$ 
15 replace every variable  $v$  in  $F$  with rename( $v$ );
16 return  $(F, X_I, \omega);$ 
17 Function rename( $v$ ):
18      $S \leftarrow \{u \in X_P \mid u \leq v\};$ 
19     if  $S = \emptyset$  then return  $v$ ;
20     return  $v - d(\max S);$ 
    
```

4.1 Proof of Correctness

Theorem 1 (Early Projection [12, 13], verbatim). *Let X and Y be sets of variables. For all functions $f: 2^X \rightarrow \mathbb{R}$ and $g: 2^Y \rightarrow \mathbb{R}$, if $x \in X \setminus Y$, then $\exists_x(f \cdot g) = (\exists_x f) \cdot g$.*

Lemma 1. *For any pseudo-Boolean function $f: 2^X \rightarrow \mathbb{R}$,*

$$(\exists_X f)(\emptyset) = \sum_{Y \subseteq X} f(Y).$$

Proof. If $X = \{x\}$, then

$$(\exists_x f)(\emptyset) = (f|_{x=1} + f|_{x=0})(\emptyset) = f|_{x=1}(\emptyset) + f|_{x=0}(\emptyset) = \sum_{Y \subseteq \{x\}} f(Y).$$

This easily extends to $|X| > 1$ by the definition of projection on sets of variables.

Proposition 2. *Let (ϕ, X_I, X_P, w) be a WMC instance. Then*

$$\left(\{[c]_0^1 \mid c \in \phi\} \cup \left\{ [x]_{w(\neg x)}^{w(x)} \mid x \in X_I \cup X_P \right\}, X_I \cup X_P, 1 \right) \quad (1)$$

is a PBP instance with the same answer.

Proof. Let $f = \prod_{c \in \phi} [c]_0^1$, and $g = \prod_{x \in X_I \cup X_P} [x]_{w(\neg x)}^{w(x)}$. Then the answer of Eq. (1) is

$$(\exists_{X_I \cup X_P} fg)(\emptyset) = \sum_{Y \subseteq X_I \cup X_P} (fg)(Y) = \sum_{Y \subseteq X_I \cup X_P} f(Y)g(Y)$$

by Lemma 1. Note that

$$f(Y) = \begin{cases} 1 & \text{if } Y \models \phi, \\ 0 & \text{otherwise,} \end{cases} \quad \text{and} \quad g(Y) = \prod_{Y \models l} w(l),$$

which means that

$$\sum_{Y \subseteq X_I \cup X_P} f(Y)g(Y) = \sum_{Y \models \phi} \prod_{Y \models l} w(l)$$

as required.

Theorem 2 (Correctness). *Algorithm 1, when given a WMC instance (ϕ, X_I, X_P, w) , returns PBP instance with the same answer, provided the following conditions are satisfied:*

1. for all indicator variables $i \in X_I$, $w(i) = w(\neg i) = 1$,
2. and either
 - (a) for all parameter variables $p \in X_P$, there is a non-empty family of literals $(l_i)_{i=1}^n$ such that
 - i. $w(\neg p) = 1$,
 - ii. $l_i \in X_I$ or $\neg l_i \in X_I$ for all $i = 1, \dots, n$,
 - iii. and $\{c \in \phi \mid p \in c \text{ or } \neg p \in c\} = \{p \vee \bigvee_{i=1}^n \neg l_i\} \cup \{l_i \vee \neg p \mid i = 1, \dots, n\}$;
 - (b) or for all parameter variables $p \in X_P$,
 - i. $w(p) + w(\neg p) = 1$,
 - ii. for any clause $c \in \phi$, $|c \cap X_P| \leq 1$,
 - iii. there is no clause $c \in \phi$ such that $\neg p \in c$,
 - iv. if $\{p\} \in \phi$, then there is no clause $c \in \phi$ such that $c \neq \{p\}$ and $p \in c$,
 - v. and for any $c, d \in \phi$ such that $c \neq d$, $p \in c$ and $p \in d$, $\bigwedge_{l \in c \setminus \{p\}} \neg l \wedge \bigwedge_{l \in d \setminus \{p\}} \neg l$ is false.

Proof. By Proposition 2,

$$\left(\{[c]_0^1 \mid c \in \phi\} \cup \left\{ [x]_{w(\neg x)}^{w(x)} \mid x \in X_I \cup X_P \right\}, X_I \cup X_P, 1 \right) \quad (2)$$

is a PBP instance with the same answer as the given WMC instance. By Definition 3, its answer is

$$\left(\exists_{X_I \cup X_P} \left(\prod_{c \in \phi} [c]_0^1 \right) \prod_{x \in X_I \cup X_P} [x]_{w(\neg x)}^{w(x)} \right) (\emptyset) \quad (3)$$

Since both Conditions 2a and 2b ensure that each clause in ϕ has at most one parameter variable, we can partition ϕ into $\phi_* := \{c \in \phi \mid \text{Vars}(c) \cap X_P = \emptyset\}$ and $\phi_p := \{c \in \phi \mid \text{Vars}(c) \cap X_P = \{p\}\}$ for all $p \in X_P$. We can then use Theorem 1 to reorder (3) into

$$\left(\exists_{X_I} \left(\prod_{x \in X_I} [x]_{w(\neg x)}^{w(x)} \right) \left(\prod_{c \in \phi_*} [c]_0^1 \right) \prod_{p \in X_P} \exists_p [p]_{w(\neg p)}^{w(p)} \prod_{c \in \phi_p} [c]_0^1 \right) (\emptyset).$$

Let us first consider how the unfinished WMC instance (F, X_I, ω) after the loop on Lines 4 to 11 differs from (2). Note that Algorithm 1 leaves each $c \in \phi_*$ unchanged, i.e., adds $[c]_0^1$ to F . We can then fix an arbitrary $p \in X_P$ and let F_p be the set of functions added to F as a replacement of ϕ_p . It is sufficient to show that

$$\omega \prod_{f \in F_p} f = \exists_p [p]_{w(\neg p)}^{w(p)} \prod_{c \in \phi_p} [c]_0^1. \quad (4)$$

Note that under Condition 2a,

$$\bigwedge_{c \in \phi_p} c \equiv p \Leftrightarrow \bigwedge_{i=1}^n l_i$$

for some family of indicator variable literals $(l_i)_{i=1}^n$. Thus,

$$\exists_p [p]_{w(\neg p)}^{w(p)} \prod_{c \in \phi_p} [c]_0^1 = \exists_p [p]_1^{w(p)} \left[p \Leftrightarrow \bigwedge_{i=1}^n l_i \right]_0^1.$$

If $w(p) = 1$, then

$$\exists_p [p]_1^{w(p)} \left[p \Leftrightarrow \bigwedge_{i=1}^n l_i \right]_0^1 = \left[p \Leftrightarrow \bigwedge_{i=1}^n l_i \right]_{0|_{p=1}}^1 + \left[p \Leftrightarrow \bigwedge_{i=1}^n l_i \right]_{0|_{p=0}}^1. \quad (5)$$

Since for any input, $\bigwedge_{i=1}^n l_i$ is either true or false, exactly one of the two summands in Eq. (5) will be equal to one, and the other will be equal to zero, and so

$$\left[p \Leftrightarrow \bigwedge_{i=1}^n l_i \right]_{0|_{p=1}}^1 + \left[p \Leftrightarrow \bigwedge_{i=1}^n l_i \right]_{0|_{p=0}}^1 = 1,$$

where 1 is a pseudo-Boolean function that always returns one. On the other side of Eq. (4), since $F_p = \emptyset$, and ω is unchanged, we get $\omega \prod_{f \in F_p} f = 1$, and so Eq. (4) is satisfied under Condition 2a when $w(p) = 1$.

If $w(p) \neq 1$, then

$$F_p = \left\{ \left[\bigwedge_{i=1}^n l_i \right]_1^{w(p)} \right\},$$

and $\omega = 1$, and so we want to show that

$$\left[\bigwedge_{i=1}^n l_i \right]_1^{w(p)} = \exists_p [p]_1^{w(p)} \left[p \Leftrightarrow \bigwedge_{i=1}^n l_i \right]_0^1,$$

and indeed

$$\exists_p [p]_1^{w(p)} \left[p \Leftrightarrow \bigwedge_{i=1}^n l_i \right]_0^1 = w(p) \cdot \left[\bigwedge_{i=1}^n l_i \right]_0^1 + \left[\bigwedge_{i=1}^n l_i \right]_1^0 = \left[\bigwedge_{i=1}^n l_i \right]_1^{w(p)}.$$

This finishes the proof of the correctness of the first ‘foreach’ loop under Condition 2a.

Now let us assume Condition 2b. We still want to prove Eq. (4). If $w(p) = 1$, then $F_p = \emptyset$, and $\omega = 1$, and so the left-hand side of Eq. (4) is equal to one. Then the right-hand side is

$$\exists_p [p]_0^1 \prod_{c \in \phi_p} [c]_0^1 = \exists_p \left[p \wedge \bigwedge_{c \in \phi_p} c \right]_0^1 = \exists_p [p]_0^1 = 0 + 1 = 1$$

since $p \in c$ for every clause $c \in \phi_p$.

If $w(p) \neq 1$, and $\{p\} \in \phi_p$, then, by Condition 2(b)iv, $\phi_p = \{\{p\}\}$, and Algorithm 1 produces $F_p = \emptyset$ and $\omega = w(p)$, and so

$$\exists_p [p]_{w(\neg p)}^{w(p)} [p]_0^1 = \exists_p [p]_0^{w(p)} = w(p) = \omega \prod_{f \in F_p} f.$$

The only remaining case is when $w(p) \neq 1$ and $\{p\} \notin \phi_p$. Then $\omega = 1$, and

$$F_p = \left\{ \left[\bigwedge_{l \in c \setminus \{p\}} \neg l \right]_1^{w(p)} \mid c \in \phi_p \right\},$$

so need to show that

$$\prod_{c \in \phi_p} \left[\bigwedge_{l \in c \setminus \{p\}} \neg l \right]_1^{w(p)} = \exists_p [p]_{1-w(p)}^{w(p)} \prod_{c \in \phi_p} [c]_0^1.$$

We can rearrange the right-hand side as

$$\begin{aligned}
\exists_p [p]_{1-w(p)}^w \prod_{c \in \phi_p} [c]_0^1 &= \exists_p [p]_{1-w(p)}^w \left[\bigwedge_{c \in \phi_p} c \right]_0^1 = \exists_p [p]_{1-w(p)}^w \left[p \vee \bigwedge_{c \in \phi_p} c \setminus \{p\} \right]_0^1 \\
&= w(p) + (1 - w(p)) \left[\bigwedge_{c \in \phi_p} c \setminus \{p\} \right]_0^1 = w(p) + \left[\bigwedge_{c \in \phi_p} c \setminus \{p\} \right]_0^{1-w(p)} \\
&= \left[\bigwedge_{c \in \phi_p} c \setminus \{p\} \right]_{w(p)}^1 = \left[\neg \bigwedge_{c \in \phi_p} c \setminus \{p\} \right]_1^{w(p)} = \left[\bigvee_{c \in \phi_p} \neg(c \setminus \{p\}) \right]_1^{w(p)} \\
&= \left[\bigvee_{c \in \phi_p} \neg \bigvee_{l \in c \setminus \{p\}} l \right]_1^{w(p)} = \left[\bigvee_{c \in \phi_p} \bigwedge_{l \in c \setminus \{p\}} \neg l \right]_1^{w(p)}.
\end{aligned}$$

By Condition 2(b)v, $\bigwedge_{l \in c \setminus \{p\}} \neg l$ can be true for at most one $c \in \phi_p$, and so

$$\left[\bigvee_{c \in \phi_p} \bigwedge_{l \in c \setminus \{p\}} \neg l \right]_1^{w(p)} = \prod_{c \in \phi_p} \left[\bigwedge_{l \in c \setminus \{p\}} \neg l \right]_1^{w(p)}$$

which is exactly what we needed to show. This ends the proof that the first loop of Algorithm 1 preserves the answer under both Condition 2a and Condition 2b. Finally, the loop on Lines 12 to 14 of Algorithm 1 replaces $[v]_1^p [\neg v]_1^q$ with $[v]_q^p$ (for some $v \in X_I$ and $p, q \in \mathbb{R}$), but, of course,

$$[v]_1^p [\neg v]_1^q = [v]_1^p [v]_q^1 = [v]_q^p,$$

i.e., the answer is unchanged.

5 Experimental Results

Notes.

- The experiments were run on a computing cluster with Intel Xeon E5-2630, Intel Xeon E7-4820, and Intel Xeon Gold 6138 processors with a 1000 s time-out separately on both encoding and inference, and a 32 GiB memory limit.³
- Along with DPMC⁴, in our experiments we also include WMC algorithms used in the papers that introduce each encoding: Ace for cd05, cd06, and d02; Cachet⁵ [15] for sbk05; and c2d⁶ [10] with query-dnnf⁷ for bklm16.

³ Each instance was run on the same processor across all algorithms and encodings.

⁴ <https://github.com/vardigroup/DPMC>

⁵ <https://cs.rochester.edu/u/kautz/Cachet/>

⁶ <http://reasoning.cs.ucla.edu/c2d/>

⁷ <http://www.cril.univ-artois.fr/kc/d-DNNF-reasoner.html>

- DPMC was run in *greedy mode*. This mode (which was not part of the original paper [13]) optimises the order in which pseudo-Boolean functions are multiplied by prioritising functions with small representations.
- Indeed, the numbers of variables in **bklm16++** and **d02++** encodings are equal whenever all variables in a Bayesian network have two values (which applies to most of the benchmark instances). Nonetheless, **bklm16++** performs much faster than **d02++**.
- DPMC is run with tree decomposition based planning and algebraic decision diagram based execution—the best-performing combination in the original set of experiments [13]. We use a single iteration of **htd** [1] to generate approximately optimal tree decompositions—we found that this configuration is efficient enough to handle huge instances, and yet the width of the returned decomposition is unlikely to differ from optimal by more than one or two.
- We also record numbers of variables in each encoding before and after our additional transformation.
- Experimental research questions:
 - Can parameter variable elimination improve inference speed?
 - How does DPMC combined with encodings without (and with) parameter variables compare with other WMC algorithms and other encodings?
 - Which instances is our approach particularly successful on (compared to other algorithms and encodings and to the same encoding before our transformation)?
 - What proportion of variables is typically removed?
 - Do some encodings benefit from this transformation more than others?
- TODO: How many times am I faster than old **cd06** and new **bklm16**? Perhaps add other numbers as well.
- To perform Bayesian network inference with DPMC (or with any other WMC algorithm not based on compilation such as **Cachet**), one needs to select a probability to compute [13, 15]. If a network comes with an evidence file, we compute the probability of this evidence. Otherwise, let X be the variable last mentioned in the Bayesian network file. If **true** is one of the values of X , then we compute $\Pr(X = \text{true})$, otherwise we choose the first-mentioned value of X .

For our experiments we use Bayesian networks available with **Ace** and **Cachet**. We split them into the following groups:

- DQMR and
- Grid networks as described by Sang et al. [16];
- Mastermind and
- Random Blocks by Chavira et al. [6];
- other binary Bayesian networks including Plan Recognition [16], Friends and Smokers, Students and Professors [6], and **tcc4f**;
- non-binary classic Bayesian networks: **alarm**, **diabetes**, **hailfinder**, **mildew**, **munin1-4**, **pathfinder**, **pigs**, and **water**.

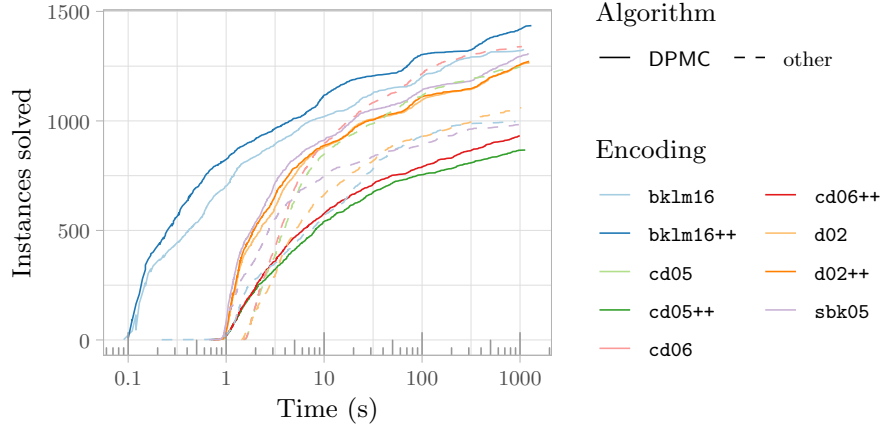


Fig. 1. Cumulative numbers of instances solved by each algorithm-encoding pair over time.

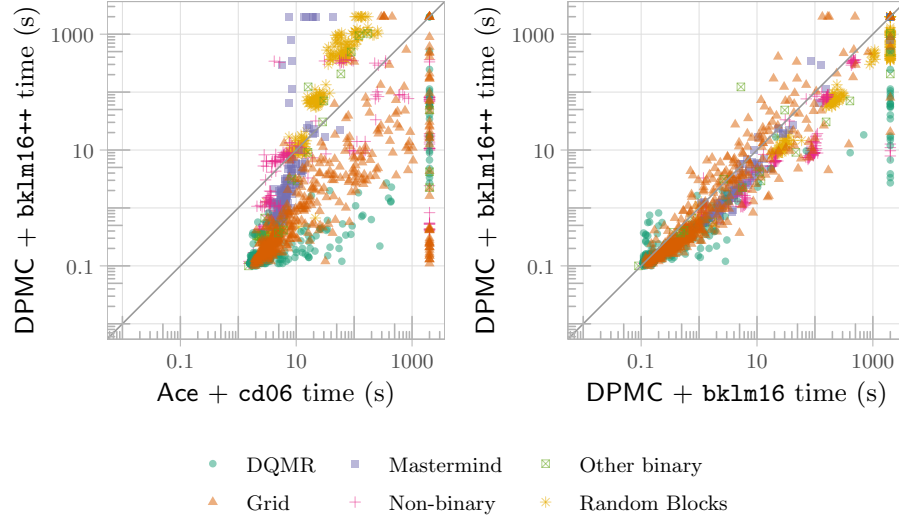


Fig. 2. An instance-by-instance comparison between DPMC+bklm16++ (the best combination according to Fig. 1) and the second and third best performing combinations: Ace + cd06 and DPMC + bklm16.

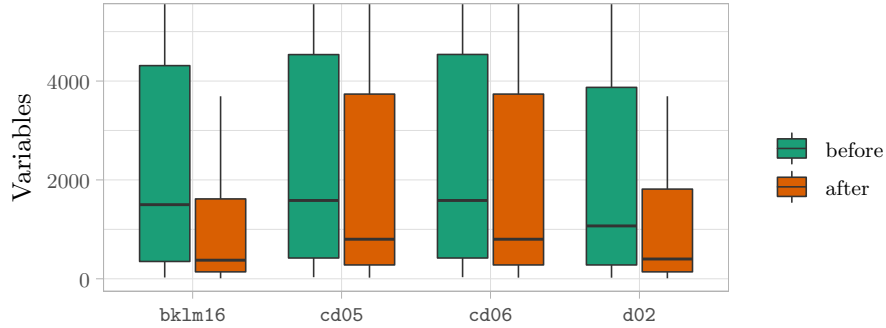


Fig. 3. Box plots of the numbers of variables in each encoding across all benchmark instances before and after applying Algorithm 1.

6 Conclusions

Notes.

- Benefits of my approach:
 - smaller primal graph, so easier to perform tree decomposition
 - Variable order is less likely to be obstructed by all the unnecessary ‘parameter’ variables.
 - There are others, but they’re not that important.

TODO

- Do I need to formally consider extending a pseudo-Boolean function to a bigger domain?
- check if each condition is actually used. Maybe turn this into a paragraph that gives an overview of the proof.
- condition 1 is necessary in both cases because we’re ignoring the weights of indicator variables (explicitly acknowledge this)
- add the cd05/cd06 correctness theorem

References

1. Abseher, M., Musliu, N., Woltran, S.: htd - A free, open-source framework for (customized) tree decompositions and beyond. In: Salvagnin, D., Lombardi, M. (eds.) *Integration of AI and OR Techniques in Constraint Programming - 14th International Conference, CPAIOR 2017, Padua, Italy, June 5-8, 2017, Proceedings*. Lecture Notes in Computer Science, vol. 10335, pp. 376–386. Springer (2017). https://doi.org/10.1007/978-3-319-59776-8_30

2. Bart, A., Koriche, F., Lagniez, J., Marquis, P.: An improved CNF encoding scheme for probabilistic inference. In: Kaminka, G.A., Fox, M., Bouquet, P., Hüllermeier, E., Dignum, V., Dignum, F., van Harmelen, F. (eds.) ECAI 2016 - 22nd European Conference on Artificial Intelligence, 29 August-2 September 2016, The Hague, The Netherlands - Including Prestigious Applications of Artificial Intelligence (PAIS 2016). *Frontiers in Artificial Intelligence and Applications*, vol. 285, pp. 613–621. IOS Press (2016). <https://doi.org/10.3233/978-1-61499-672-9-613>
3. Chavira, M., Darwiche, A.: Compiling Bayesian networks with local structure. In: Kaelbling, L.P., Saffioti, A. (eds.) IJCAI-05, Proceedings of the Nineteenth International Joint Conference on Artificial Intelligence, Edinburgh, Scotland, UK, July 30 - August 5, 2005. pp. 1306–1312. Professional Book Center (2005), <http://ijcai.org/Proceedings/05/Papers/0931.pdf>
4. Chavira, M., Darwiche, A.: Encoding CNFs to empower component analysis. In: Biere, A., Gomes, C.P. (eds.) Theory and Applications of Satisfiability Testing - SAT 2006, 9th International Conference, Seattle, WA, USA, August 12-15, 2006, Proceedings. *Lecture Notes in Computer Science*, vol. 4121, pp. 61–74. Springer (2006). https://doi.org/10.1007/11814948_9
5. Chavira, M., Darwiche, A.: On probabilistic inference by weighted model counting. *Artif. Intell.* **172**(6-7), 772–799 (2008). <https://doi.org/10.1016/j.artint.2007.11.002>
6. Chavira, M., Darwiche, A., Jaeger, M.: Compiling relational bayesian networks for exact inference. *Int. J. Approx. Reason.* **42**(1-2), 4–20 (2006). <https://doi.org/10.1016/j.ijar.2005.10.001>
7. Choi, A., Kisa, D., Darwiche, A.: Compiling probabilistic graphical models using sentential decision diagrams. In: van der Gaag, L.C. (ed.) Symbolic and Quantitative Approaches to Reasoning with Uncertainty - 12th European Conference, ECSQARU 2013, Utrecht, The Netherlands, July 8-10, 2013. Proceedings. *Lecture Notes in Computer Science*, vol. 7958, pp. 121–132. Springer (2013). https://doi.org/10.1007/978-3-642-39091-3_11
8. Darwiche, A.: On the tractable counting of theory models and its application to truth maintenance and belief revision. *J. Appl. Non Class. Logics* **11**(1-2), 11–34 (2001). <https://doi.org/10.3166/jancl.11.11-34>
9. Darwiche, A.: A logical approach to factoring belief networks. In: Fensel, D., Giunchiglia, F., McGuinness, D.L., Williams, M. (eds.) Proceedings of the Eight International Conference on Principles and Knowledge Representation and Reasoning (KR-02), Toulouse, France, April 22-25, 2002. pp. 409–420. Morgan Kaufmann (2002)
10. Darwiche, A.: New advances in compiling CNF into decomposable negation normal form. In: de Mántaras, R.L., Saitta, L. (eds.) Proceedings of the 16th European Conference on Artificial Intelligence, ECAI’2004, including Prestigious Applicants of Intelligent Systems, PAIS 2004, Valencia, Spain, August 22-27, 2004. pp. 328–332. IOS Press (2004)
11. Darwiche, A.: SDD: A new canonical representation of propositional knowledge bases. In: Walsh, T. (ed.) IJCAI 2011, Proceedings of the 22nd International Joint Conference on Artificial Intelligence, Barcelona, Catalonia, Spain, July 16-22, 2011. pp. 819–826. IJCAI/AAAI (2011). <https://doi.org/10.5591/978-1-57735-516-8/IJCAI11-143>, <https://doi.org/10.5591/978-1-57735-516-8/IJCAI11-143>
12. Dudek, J.M., Phan, V., Vardi, M.Y.: ADDMC: weighted model counting with algebraic decision diagrams. In: The Thirty-Fourth AAAI Conference on Artificial Intelligence, AAAI 2020, The Thirty-Second Innovative Applications

- of Artificial Intelligence Conference, IAAI 2020, The Tenth AAAI Symposium on Educational Advances in Artificial Intelligence, EAAI 2020, New York, NY, USA, February 7-12, 2020. pp. 1468–1476. AAAI Press (2020), <https://aaai.org/ojs/index.php/AAAI/article/view/5505>
13. Dudek, J.M., Phan, V.H.N., Vardi, M.Y.: DPMC: weighted model counting by dynamic programming on project-join trees. In: Simonis, H. (ed.) *Principles and Practice of Constraint Programming - 26th International Conference, CP 2020, Louvain-la-Neuve, Belgium, September 7-11, 2020, Proceedings. Lecture Notes in Computer Science*, vol. 12333, pp. 211–230. Springer (2020). https://doi.org/10.1007/978-3-030-58475-7_13
 14. Oztok, U., Darwiche, A.: A top-down compiler for sentential decision diagrams. In: Yang, Q., Wooldridge, M.J. (eds.) *Proceedings of the Twenty-Fourth International Joint Conference on Artificial Intelligence, IJCAI 2015, Buenos Aires, Argentina, July 25-31, 2015*. pp. 3141–3148. AAAI Press (2015), <http://ijcai.org/Abstract/15/443>
 15. Sang, T., Bacchus, F., Beame, P., Kautz, H.A., Pitassi, T.: Combining component caching and clause learning for effective model counting. In: *SAT 2004 - The Seventh International Conference on Theory and Applications of Satisfiability Testing*, 10-13 May 2004, Vancouver, BC, Canada, Online Proceedings (2004), <http://www.satisfiability.org/SAT04/programme/21.pdf>
 16. Sang, T., Beame, P., Kautz, H.A.: Performing Bayesian inference by weighted model counting. In: Veloso, M.M., Kambhampati, S. (eds.) *Proceedings, The Twentieth National Conference on Artificial Intelligence and the Seventeenth Innovative Applications of Artificial Intelligence Conference*, July 9-13, 2005, Pittsburgh, Pennsylvania, USA. pp. 475–482. AAAI Press / The MIT Press (2005), <http://www.aaai.org/Library/AAAI/2005/aaai05-075.php>
 17. Shen, Y.: *Modeling, Learning and Reasoning with Structured Bayesian Networks*. Ph.D. thesis, UCLA (2020)